

docker

4PM



Docker overview

1



Images

2



Containers

3

6

Docker Swarm

5

Docker Compose

4


Networking &  
Volume




# What is a docker ?

image


- Docker is containerization platform that enables developer to build, test and deploy the application easily and reliably
- Docker host runs multiple containers maintaining the isolation between the containers



machine




machine



machine


# Container

- Allows developers to create and deploy applications faster and more securely
- A container is a standard unit which provides single point of service → *application + libraries*
- Involves encapsulating or packaging up software code and all its dependencies so that it can run uniformly and consistently on any infrastructure
- It is a operating system virtualization



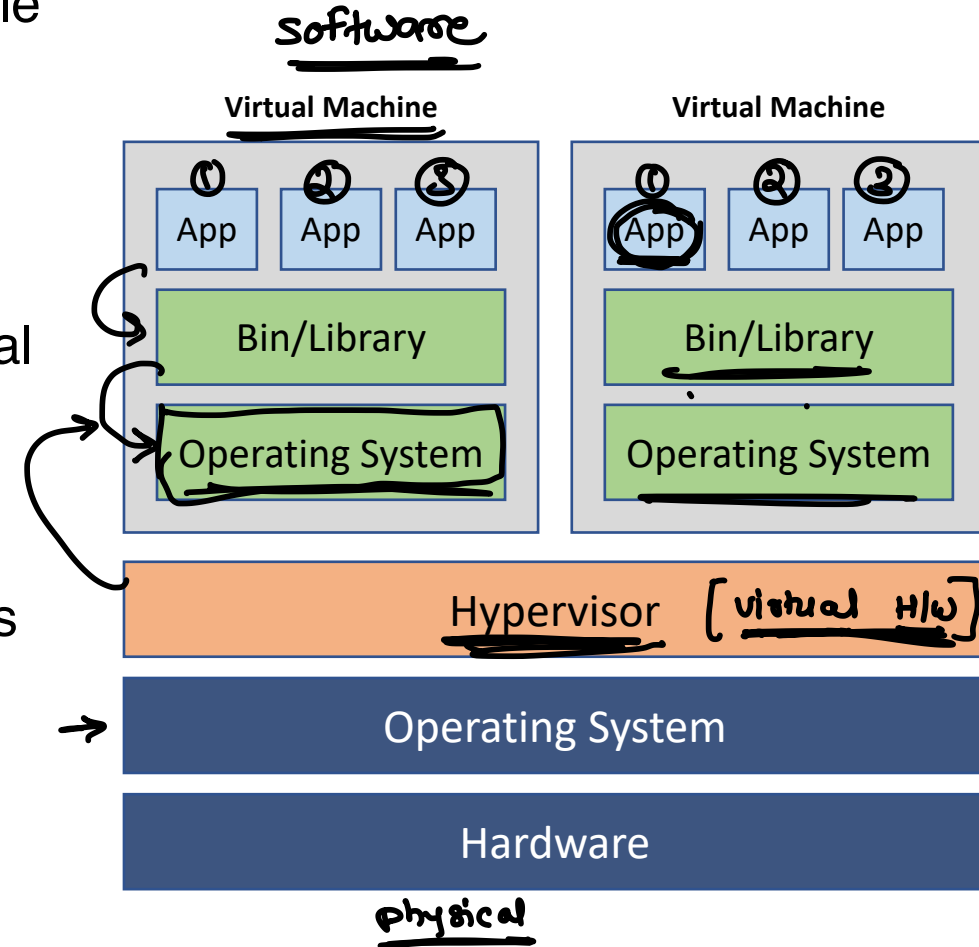
# Traditional Deployment

- Early on, organizations ran applications on physical servers
- There was no way to define resource boundaries for applications in a physical server, and this caused resource allocation issues
- For example, if multiple applications run on a physical server, there can be instances where one application would take up most of the resources, and as a result, the other applications would underperform
- A solution for this would be to run each application on a different physical server
- But this did not scale as resources were underutilized, and it was expensive for organizations to maintain many physical servers



# Virtualized Deployment


- It allows you to run multiple Virtual Machines (VMs) on a single physical server's CPU
- Virtualization allows applications to be isolated between VMs and provides a level of security as the information of one application cannot be freely accessed by another application
- Virtualization allows better utilization of resources in a physical server and allows better scalability because
  - an application can be added or updated easily
  - reduces hardware costs
- With virtualization you can present a set of physical resources as a cluster of disposable virtual machines
- Each VM is a full machine running all the components, including its own operating system, on top of the virtualized hardware




# Container deployment

Linux

- Containers are similar to VMs, but they have relaxed isolation properties to share the Operating System (OS) among the applications
- Therefore, containers are considered lightweight
- Similar to a VM, a container has its own filesystem, CPU, memory, process space, and more
- As they are decoupled from the underlying infrastructure, they are portable across clouds and OS distributions



win/macos



# Containerization vs Virtualization

immutable  
read only

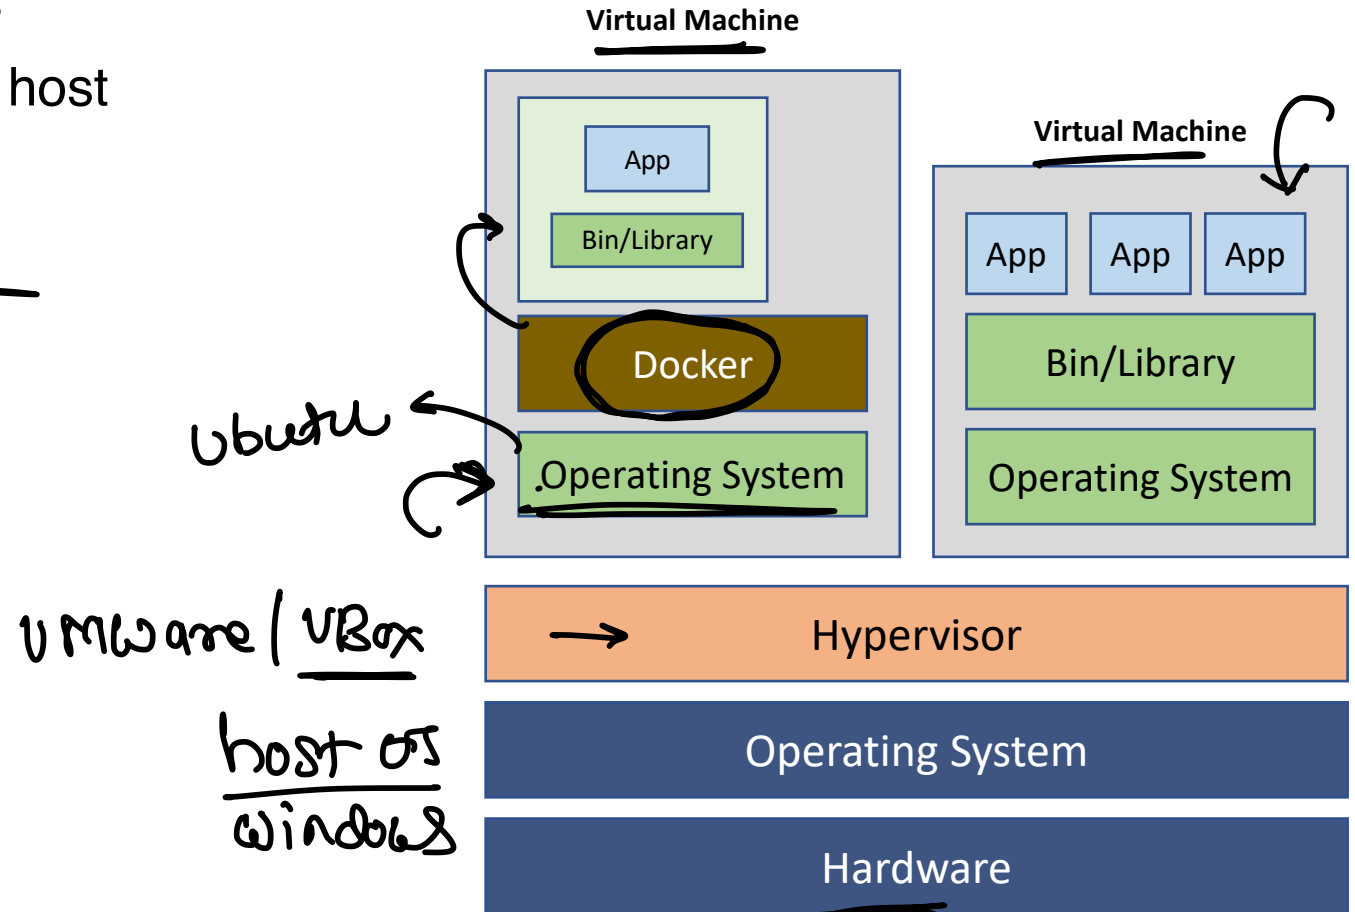

Virtual Machine	Container
<u>Hardware level virtualization</u>	<u>OS virtualization</u>
<u>Heavyweight (bigger in size)</u>	<u>Lightweight (smaller in size)</u>
<u>Slow provisioning</u>	<u>Real-time and fast provisioning</u>
<u>Limited Performance</u>	<u>Native performance</u>
<u>Fully isolated</u>	<u>Process-level isolation</u>
<u>More secure</u>	<u>Less secure</u>
<u>Each VM has separate OS</u>	<u>Each container can share OS resources</u>
<u>Boots in minutes</u>	<u>Boots in seconds</u>
<u>Pre-configured VMs are difficult to find and manage</u>	<u>Pre-built containers are readily available</u>
<u>Can be easily moved to new OS</u>	<u>Containers are destroyed and recreated</u>
<u>Creating VM takes longer time</u>	<u>Containers can be created in seconds</u>

→ virtualize HW



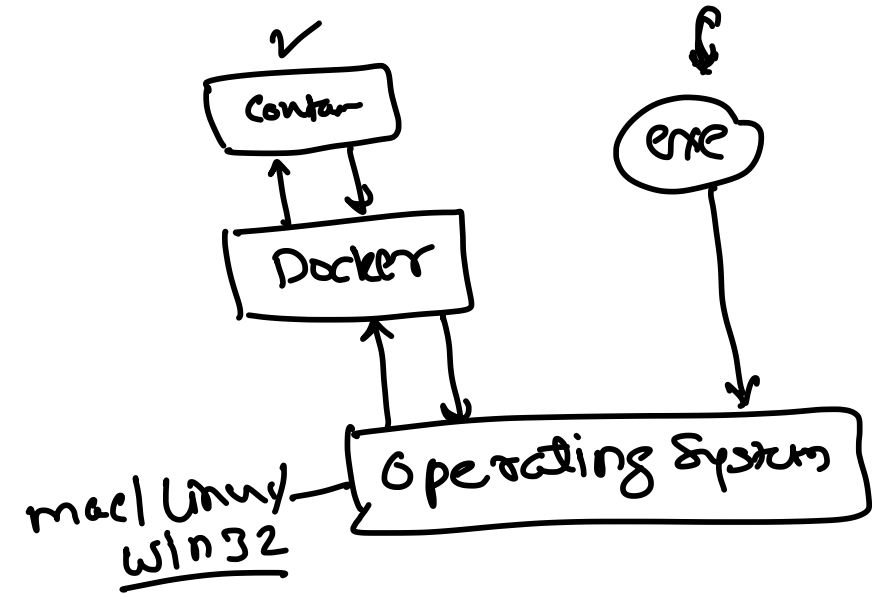
# Containerization and Virtualization

- Containers can run inside virtual machines
- In which case, the a physical machine can host VM that may house Docker containers
- This is preferred in the cloud environment



# Why Docker?

- It is an easy way to create application deployable packages
- Developer can create ready-to-run containerized applications
- It provides consistent computing environment
- It works equally well in on-prem as well as cloud environments
- It is light weight compared to VM



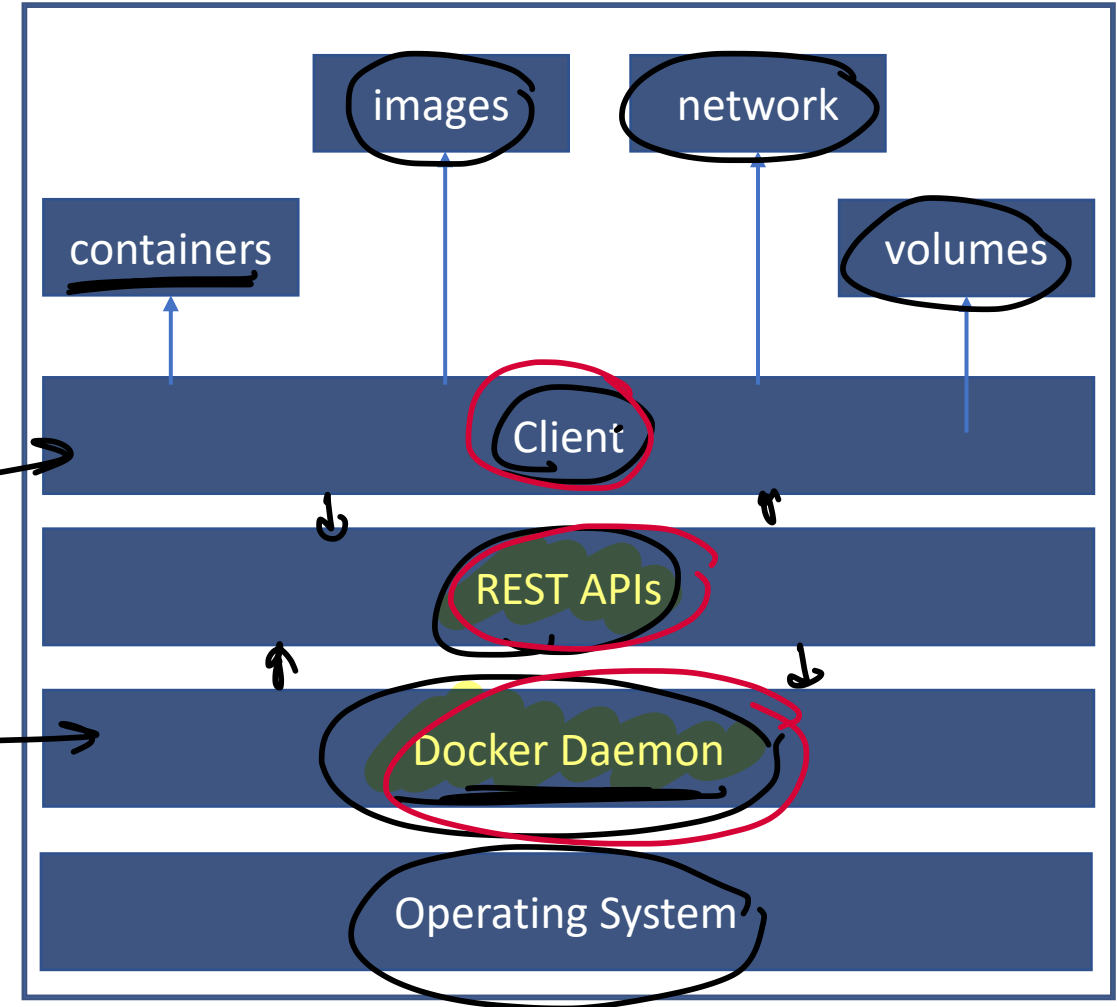
# Little history about Docker

- Docker Inc, started by Solomon Hykes, is behind the docker tool
- Docker Inc started as PaaS provider called as dotCloud
- In 2013, the dotCloud became Docker Inc
- Docker Inc was using Linux Containers (LXC) before version 0.9
- After 0.9 (2014), Docker replaced LXC with its own library libcontainer which is developed in Go programming language
- Its not the only solution for containerization
  - "FreeBSD Jails", launched in 2000
  - LXD is next generation system container manager build on top of LXC and REST APIs
  - Google has its own open source container technology Imctfy (Let Me Contain That For You)
  - Rkt is another option for running containers

# Docker Architecture

- Docker daemon (dockerd)
  - Continuous running process (server)
  - Manages the containers
- REST APIs
  - Used to communicate with docker daemon
- Client (docker)
  - Provides command line interface
  - Used to perform all the tasks

sudo systemctl start docker



# libcontainer [Go]

- Docker has replaced LXC by libcontainer, which is used to manage the containers

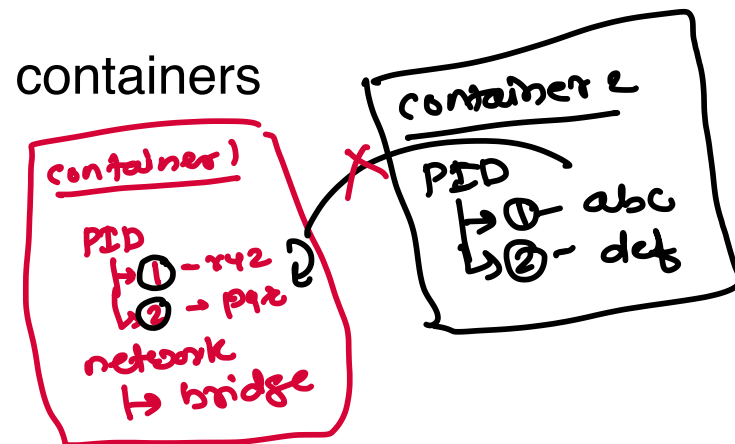
- Libcontainer uses

## ✓ Namespaces

- Creates isolated workspace which limits what container can see
- Provides a layer of isolation to the container
- Each container runs in a separate namespace
- Processes running in a namespace can interact with other processes or use resources which are the part of the same namespace
- E.g. process ID, network, IPC, Filesystem

## ✓ Control Groups (cgroups)

- Used to share the available resources to the containers
- It optionally enforces limits and constraints on resource usage
- It limits how much a container can use
- E.g. CPU, Disk space, memory



1 GB / 4 GB  
50 MB / 512 GB

# libcontainer

- ✓ Union File System (UnionFS)
  - It uses layers
  - It is a lightweight and very fast FS
  - Docker uses different variants of UnionFS
    - ✓ ▪ Aufs (advanced multi-layered unification filesystem)
    - ✓ ▪ Btrfs (B-Tree FS)
    - ✓ ▪ VFS (Virtual FS)
    - ✓ ▪ Devicemapper

# Docker Objects

---

- ✓ Images: read only template with instructions for creating docker containers
- ✓ Container: running instance of a docker image
- ✓ Network: network interface used to connect the containers to each other or external networks
- ✓ Volumes: used to persist the data generated by and used by the containers
- ✓ Registry: private or public collection of docker images
- Service: used to deploy application in a docker multi node cluster


# Docker Images





# Docker Image [class]

- Read only template which has instructions for running docker containers
- In order to run a container, developer first need to package the application along with its dependencies in the form of a docker image
- It is highly portable and can be shared over network, stored and updated → D.R. ← contains  
dockerfile
- Docker provides public or private registry which contains collection of pre-built images → Hub
- If the image you are looking for is not available publicly, you can create your image using Dockerfile
- One image can be based on another image



# Demo



---

- Tasks
  - ✓ ■ Search images on docker hub ✓
  - ✓ ■ Download/pull image ← ✓
  - ✓ ■ Get information of an image ✓
  - ✓ ■ List all pulled images ✓
    - Remove an image

# Docker Container

# Container




- It is a running aspect of docker image
- Contains one or more running processes
- It is a self-contained environment
- It wraps up an application into its own isolated box (application running inside a container has no knowledge of any other applications or processes that exist outside the container)
- A container can not modify the image from which it is created
- It consists of
  - ✓ Your application code
  - ✓ Dependencies
  - ✓ Networking
  - ✓ Volumes



# Demo

## ■ Tasks

- ✓ List the running containers on the host
- ✓ Create a container
- ✓ Start a created/stopped container
- ✓ Run a container
- ✓ Stop a container
- ✓ Restart a container



- ↳ Docker managed PS

# Attaching a container



- There are two ways to attach to a container

## Attach

- Used to attach the container
- Uses only one input and output stream
- Task
  - Attach to a running container

## Exec

- Mainly it is used for running a command inside a container
- Task
  - Execute a command inside container



# Hostname and name of container

- To check the host name
  - Go inside the container
  - Check the hostname by using a command hostname
- Docker uses the first 12 characters of container id as hostname
- Docker automatically generates a name at random for each container

- Task

- Change the hostname of a container
- Chane the container name of a container

= --hostname


= --name

> run



# Publishing port on container

- Publishing a port is required to give an external access to your application
- Port can be published only at the time of creating a container
- You can not update the port configuration on running container
- Task
  - Run a httpd container with port 8080 published, to access apache externally



# Container information

---

- Docker top command      ← top processes
- Docker stats command      ← summary
- Docker inspect command      ← information

# Deleting container

---

- To delete container use rm command
- To delete the container automatically use `--rm`

# Docker Network

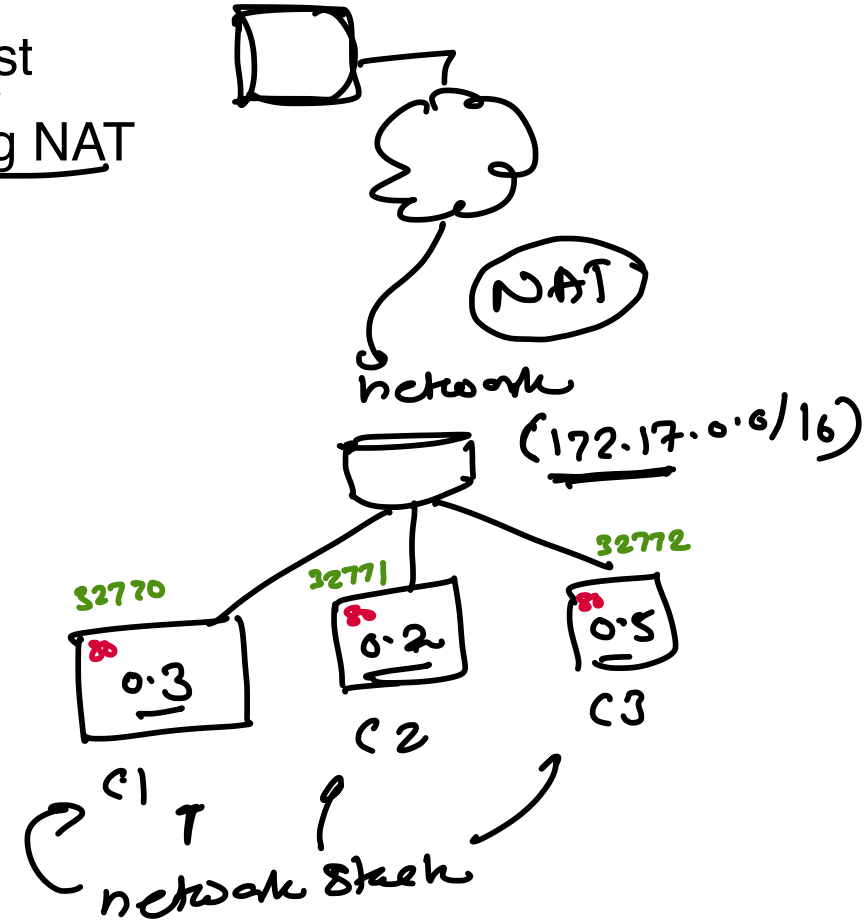
# Overview

---

- By default docker creates following networks on the host
  - ✓ ■ Bridge
  - ✓ ■ Host – only
  - ✓ ■ None
- Task
  - Check the networks on the host machine
  - Get more information of any network


# Bridge network

- Containers run on a separate network stacks, internal to docker host
- All of the containers share the external IP of the host machine using NAT
- Docker by default puts new container on bridge network
- Task
  - Get information about the bridge network
  - Run two containers on bridge network with same port published



# Host network

- Containers behave just as any other process running in the docker host
- Host network adds the containers on the host's network stack
- There will be no isolation between the host machine and the container
- Does not perform any operation on incoming traffic (NAT)
- Task
  - Run a container on host network and verify the IP address
  - Run two containers on host network with same port published



# Modify network settings on container

---

- Docker allows to modify the network settings without the need to restart the container
- Tasks
  - Start a container on none network
  - Disconnect the none network
  - Connect to bridge network



# Custom network

- Docker network command can be used to create custom networks
- To create a custom network we have to use a driver → *bridge*
- If driver is not mentioned then docker uses bridge by default
- We can create as many networks as we need
- Tasks
  - Create a custom bridge network
  - Check the network interfaces on the docker host
  - Run a container using the newly created network

# Remove the network

---

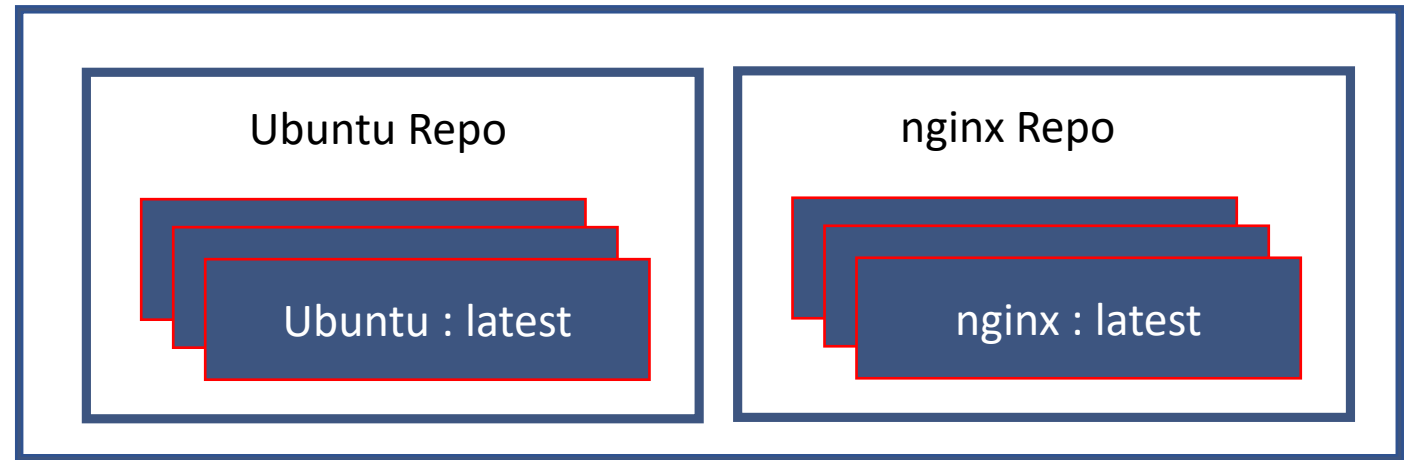
- Default networks can not be removed
- Active networks can not be removed
- Tasks
  - Create a custom network
  - Remove that custom network
- Prune command can be used to remove all unused networks

# Docker Images (Advanced)

# Docker Image

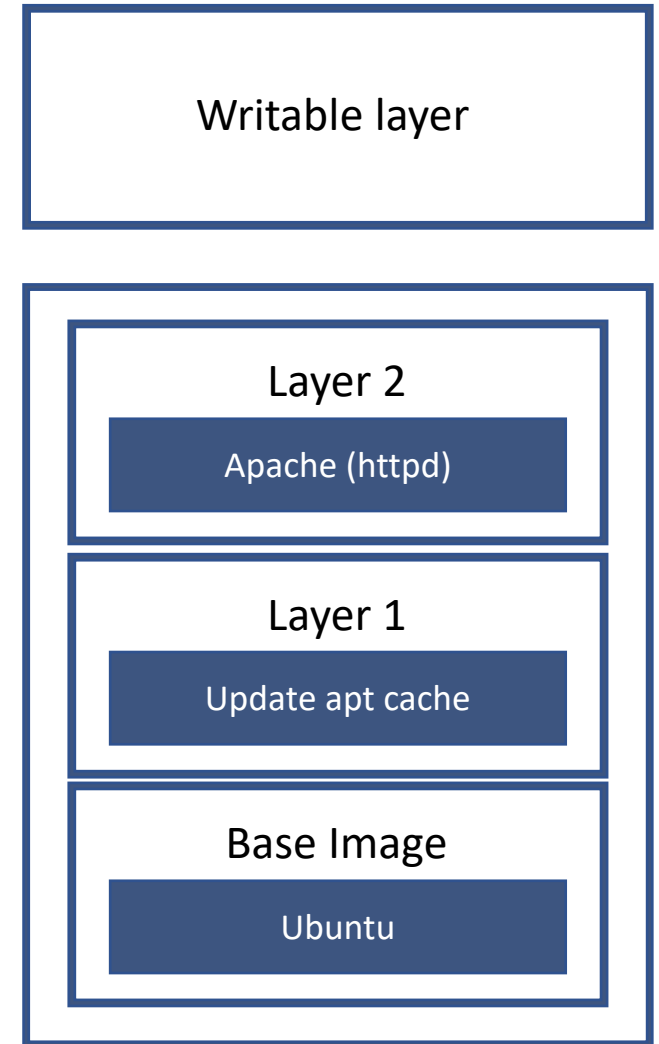
- Read-only instructions to run the containers
- It is made up of different layers
- Repositories hold images
- Docker registry stores repositories
- To create a custom image
  - Commit the running container
  - Use a Dockerfile
- Task
  - Create a container
  - Create a directory and a file within it
  - Commit the container to create a new image

## Docker Registry Server



# Layered File System

- Docker images are made of layered FS
- Docker uses UnionFS for implementing the layered docker images
- Any update on the image adds a new layer
- All changes made to the running container are written inside a writable layer



# Dockerfile

---

- The Dockerfile contains a series of instructions paired with arguments
- Each instruction should be in upper-case and be followed by an argument
- Instructions are processed from top to bottom
- Each instruction adds a new layer to the image and then commits the image
- Upon running, changes made by an instruction make it to the container

# Dockerfile instructions

---

- FROM
- ENV
- RUN
- CMD
- EXPOSE
- WORKDIR
- ADD
- COPY
- LABEL
- MAINTAINER
- ENTRYPOINT

# Sharing docker images

---

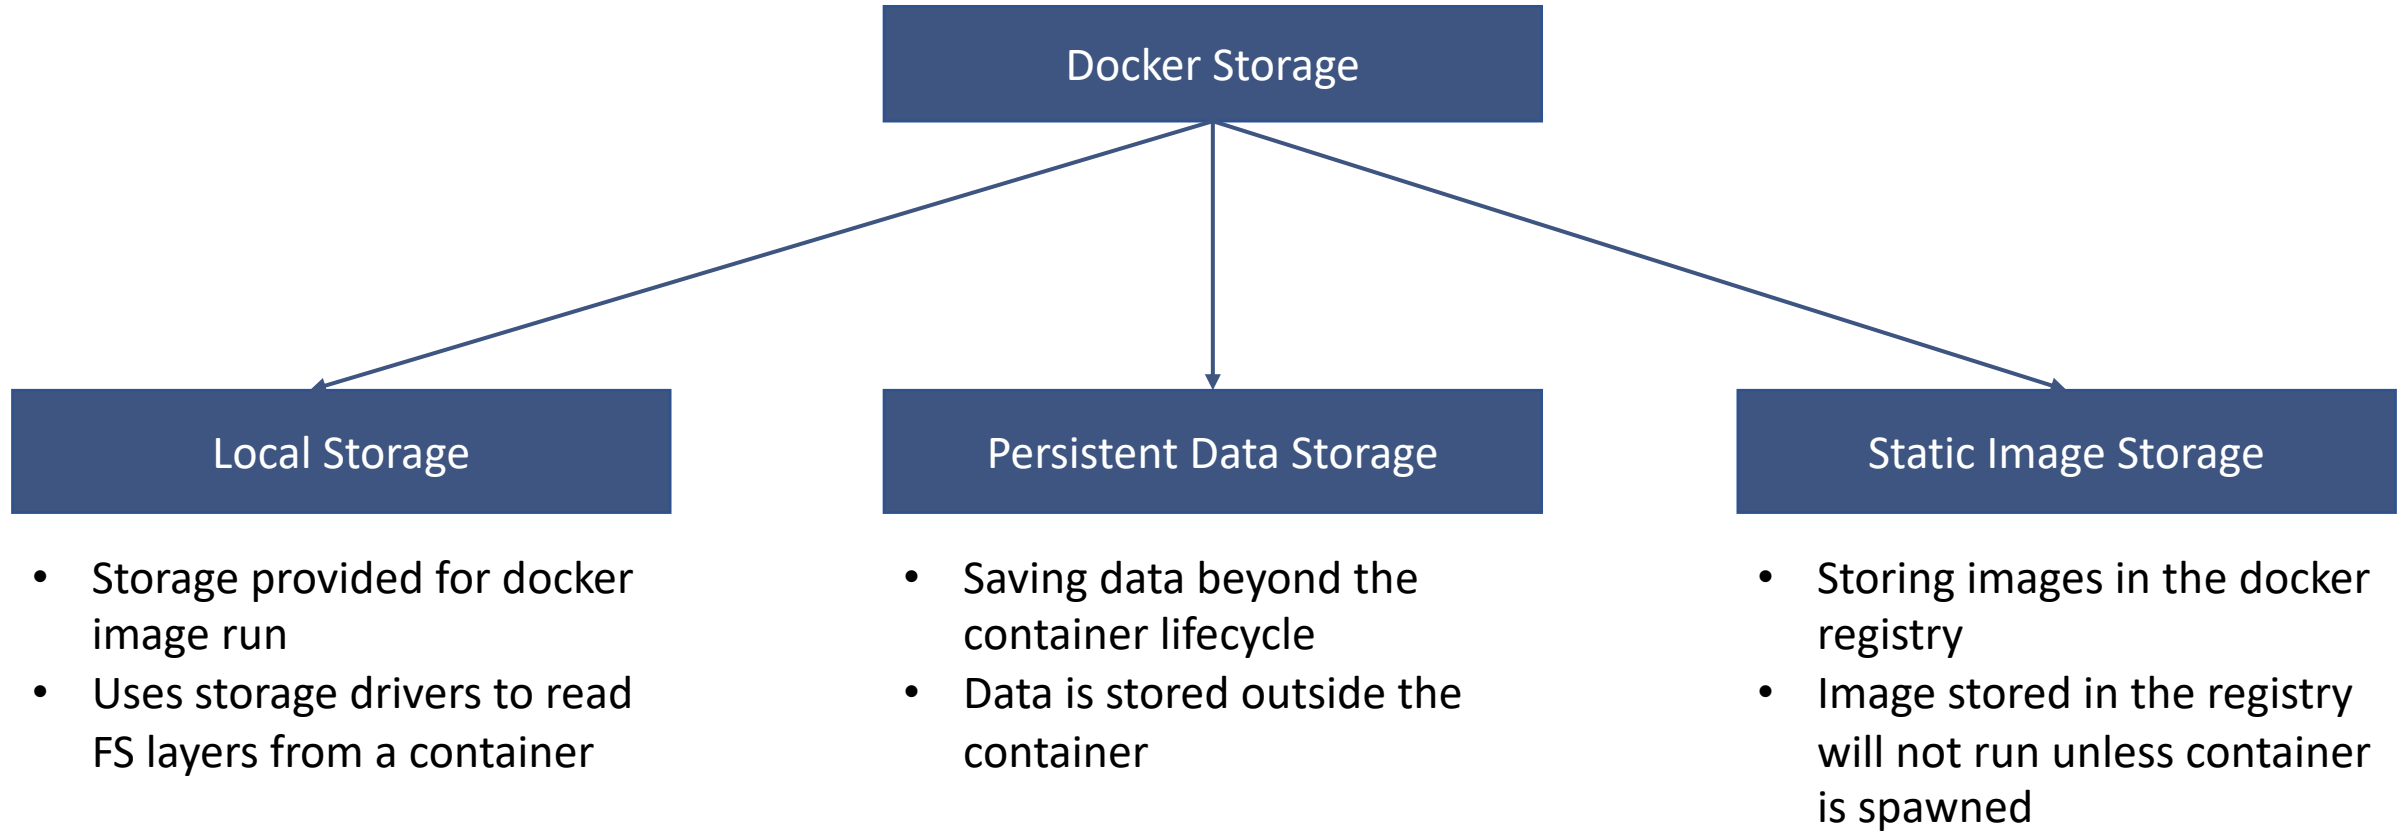
- Docker provides public docker hub to share the images
- Task
  - Create an image using Dockerfile
  - Push the image to dockerhub



# Docker Volume

# Overview

---



# Storage drivers

---

- Docker supports several different storage drivers
- E.g.
  - Overlay2
    - preferred storage driver, for all currently supported Linux distributions,
    - Requires no extra configuration
  - Aufs
    - Preferred storage driver for Docker 18.06 and older, when running on Ubuntu 14.04 on kernel 3.13 which has no support for overlay2
  - Devicemapper
    - is supported, but requires direct-lvm for production environments
  - Btrfs and zfs
    - Used if they are the backing filesystem (snapshots)
  - Vfs
    - Intended for testing purposes
- Tasks
  - Get the docker disk usage
  - Get the current storage driver configured


# Local Storage

---

- Size taken by container
  - Size: data on the writable layer
  - Virtual Size: read-only image data + writable layer size
- Multiple containers share the image hence the image size will be shared
- Tasks
  - Create a container
  - Get the size information

# Persistent Storage

- Downsides of using local storage for containers
  - Data does not persist when container is removed
  - Writable layer is tightly coupled to the host machine
- Volume provides persistent storage
- Allows to share the data among containers
- Can be managed using the docker CLI commands
- NOTE: Volume does not increase the size of container using it



# Persistent Storage

---

- Volumes
  - Stored in the docker managed FS of the host
  - Supports the use of Volume Drivers
- BindMounts
  - Stored anywhere in the host
  - E.g. you can mount a local directory with the container to share the contents
- Tmpfs Mounts
  - Temporary and stored in the host's memory
  - When the container stops, the tmpfs mount is removed
  - If the container is committed then tmpfs is not saved
  - Available only with docker on linux